

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv("emails.csv")
```

```
In [3]: df.shape
```

```
Out[3]: (5172, 3002)
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey	jay	valued	lay	infrastru
0	Email 1	0	0	1	0	0	0	2	0	0	...	0	0	0	0	
1	Email 2	8	13	24	6	6	2	102	1	27	...	0	0	0	0	
2	Email 3	0	0	1	0	0	0	8	0	0	...	0	0	0	0	
3	Email 4	0	5	22	0	5	1	51	2	10	...	0	0	0	0	
4	Email 5	7	6	17	1	5	2	57	0	9	...	0	0	0	0	

5 rows × 3002 columns



```
In [5]: #input data
x = df.drop(['Email No.', 'Prediction'], axis = 1)
#output data
y = df['Prediction']
```

```
In [6]: x.shape
```

```
Out[6]: (5172, 3000)
```


```
In [7]: x.dtypes
```

```
Out[7]: the                int64
to                int64
ect                int64
and                int64
for                int64
...
infrastructure     int64
military           int64
allowing           int64
ff                 int64
dry                int64
Length: 3000, dtype: object
```

```
In [8]: set(x.dtypes)
```

```
Out[8]: {dtype('int64')}
```

```
In [9]: import seaborn as sns  
sns.countplot(x = y);
```



```
In [10]: y.value_counts()
```

```
Out[10]: Prediction  
0      3672  
1      1500  
Name: count, dtype: int64
```

```
In [11]: #Feature Scaling  
from sklearn.preprocessing import MinMaxScaler  
scaler = MinMaxScaler()  
x_scaled = scaler.fit_transform(x)
```

```
In [12]: x_scaled
```

```
Out[12]: array([[0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
                [0.03809524, 0.09848485, 0.06705539, ..., 0.          , 0.00877193,
                0.          ],
                [0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
                ...,
                [0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
                [0.00952381, 0.0530303 , 0.          , ..., 0.          , 0.00877193,
                0.          ],
                [0.1047619 , 0.18181818, 0.01166181, ..., 0.          , 0.          ,
                0.          ]])
```

```
In [13]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(
x_scaled, y, random_state=0, test_size=0.25)
```

```
In [14]: x_scaled.shape
```

```
Out[14]: (5172, 3000)
```

```
In [15]: x_train.shape
```

```
Out[15]: (3879, 3000)
```

```
In [16]: x_test.shape
```

```
Out[16]: (1293, 3000)
```

```
In [17]: #import the class
from sklearn.neighbors import KNeighborsClassifier
```

```
In [18]: #Create the object
knn = KNeighborsClassifier(n_neighbors=5)
```

```
In [19]: #Train the algorithm
knn.fit(x_train, y_train)
```

```
Out[19]: 

▾ KNeighborsClassifier
  KNeighborsClassifier()



```

```
In [20]: #Predict on test data
y_pred = knn.predict(x_test)
```

```
In [21]: # Import the evaluation metrics
from sklearn.metrics import ConfusionMatrixDisplay, accuracy_score
from sklearn.metrics import classification_report
```

In [22]: `ConfusionMatrixDisplay.from_predictions(y_test, y_pred)`

Out[22]: `<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2c638ef7bd0>`



In [23]: `y_test.value_counts()`

Out[23]: Prediction  
 0 929  
 1 364  
 Name: count, dtype: int64

In [24]: `accuracy_score(y_test,y_pred)`

Out[24]: 0.871616395978345

In [25]: `print(classification_report(y_test,y_pred))`

	precision	recall	f1-score	support
0	0.98	0.84	0.90	929
1	0.70	0.95	0.81	364
accuracy			0.87	1293
macro avg	0.84	0.89	0.85	1293
weighted avg	0.90	0.87	0.88	1293

```
In [26]: import numpy as np
import matplotlib.pyplot as plt
```

```
In [27]: error = []
for k in range(1,41):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(x_train,y_train)
    pred = knn.predict(x_test)
    error.append(np.mean(pred != y_test))
```

```
In [28]: error
```

```
Out[28]: [0.10827532869296211,
0.10982211910286156,
0.12296983758700696,
0.11523588553750967,
0.12838360402165508,
0.1214230471771075,
0.15158546017014696,
0.14849187935034802,
0.17246713070378963,
0.16705336426914152,
0.1871616395978345,
0.18329466357308585,
0.21500386697602475,
0.21345707656612528,
0.22815158546017014,
0.2266047950502707,
0.23588553750966745,
0.23356535189481825,
0.2459396751740139,
0.24361948955916474,
0.2559938128383604,
0.2552204176334107,
0.2699149265274555,
0.2691415313225058,
0.2822892498066512,
0.28306264501160094,
0.2954369682907966,
0.2923433874709977,
0.3039443155452436,
0.300077339520495,
0.30549110595514306,
0.30549110595514306,
0.31245166279969067,
0.31245166279969067,
0.3194122196442382,
0.317092034029389,
0.32637277648878577,
0.32559938128383603,
0.33410672853828305,
0.3325599381283836]
```

```
In [29]: knn = KNeighborsClassifier(n_neighbors=1)
```

```
In [30]: knn.fit(x_train, y_train)
```

```
Out[30]: 

▼



KNeighborsClassifier



KNeighborsClassifier(n_neighbors=1)


```

```
In [31]: y_pred = knn.predict(x_test)
```

```
In [32]: accuracy_score(y_test, y_pred)
```

```
Out[32]: 0.8917246713070379
```

```
In [33]: from sklearn.svm import SVC
```

```
In [34]: svm = SVC(kernel = 'poly')
```

```
In [35]: svm.fit(x_train, y_train)
```

```
Out[35]: 

▼



SVC



SVC(kernel='poly')


```

```
In [36]: y_pred = svm.predict(x_test)
```

```
In [37]: accuracy_score(y_test, y_pred)
```

```
Out[37]: 0.7548337200309359
```

```
In [38]: #Linear : 0.9767981438515081  
#RBF : 0.9450889404485692  
#Polynomial : 0.7548337200309359
```

```
In [ ]:
```